

# Generative Adversarial Networks (GAN)


Michael(Mike) Erlihson, Ph.D.

- GANs Applications
- So how do GANs work
- GAN Loss Function
- GANs types and architectures
- GANs Shortcomings

- Type of neural network architecture that allow neural networks to generate data.
- Recently GANs became one of the hottest subfields in deep learning
- Generating fuzzy images of digits to photorealistic images of faces

# What GANs are able to generate?

< itc >



# More advanced GAN capabilities:

< itc >

Zebras ↔ Horses



zebra → horse



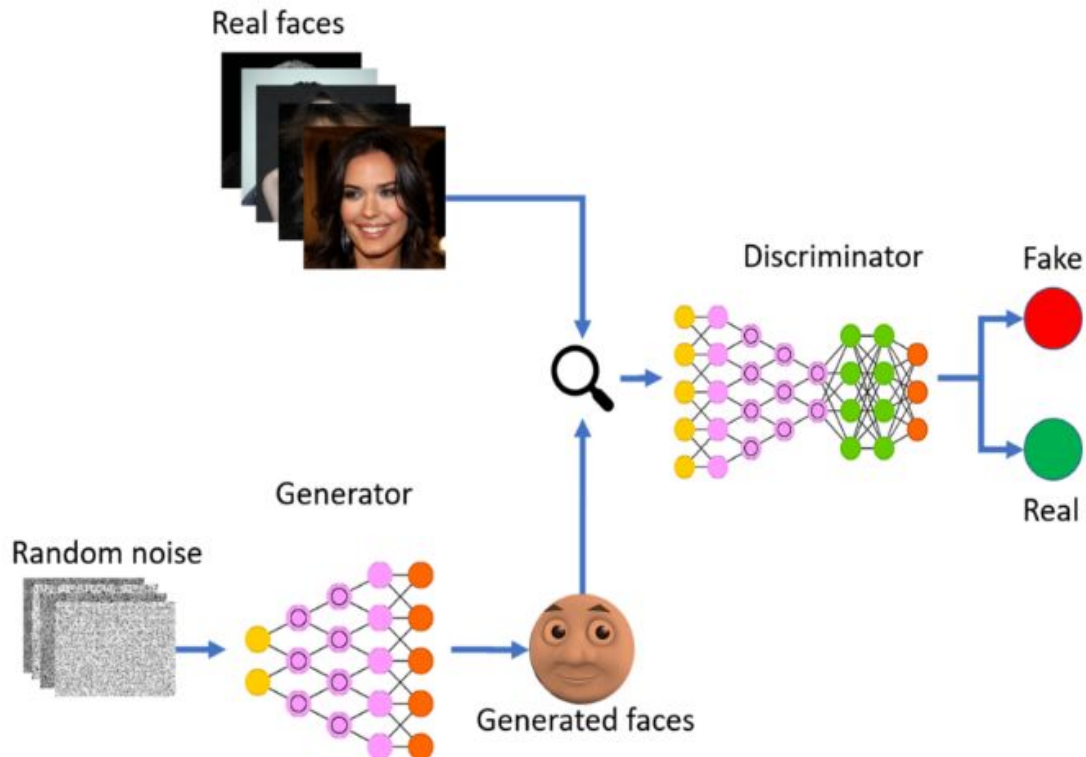
horse → zebra



- Learn dataset probability distribution via training two neural networks against each other.
- Solve two problems simultaneously:
  - **Discrimination** - distinguish between real and fake images)
  - **“Realistic” fake data generation** - construct samples similar to real
- These tasks are complete opposites - what would occur if we split them into different models: **Generator vs. Discriminator?**

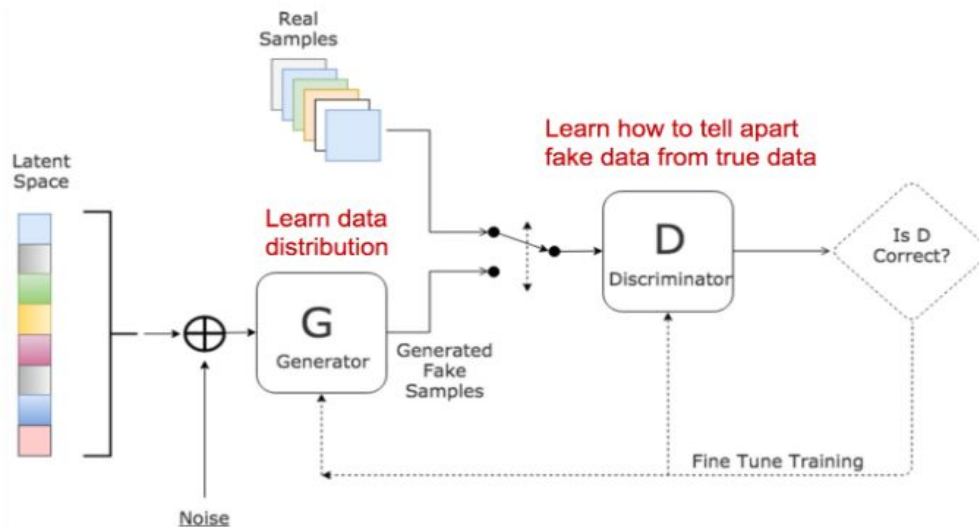
# GAN Architecture: Basic Scheme

< itc >



# GAN Architecture: An essence

- **Generator**(painting forger): tries to create images looking very similar to dataset
- **Discriminator**(police)- tries to detect whether generated images were fake or not.





# GAN Architecture: Training

< itc >



Beginning of training

Later training stage

- **Generator**(forger) is improving in creating fakes, while **discriminator**(police) keep getting better at detecting fakes.
- These two models keep competing with each other during the training process
- **Goal:** produce two networks that are each as good as they can be: we don't end up with a “winner”
- Finally the generator “learns” to create images “indistinguishable” from the real dataset

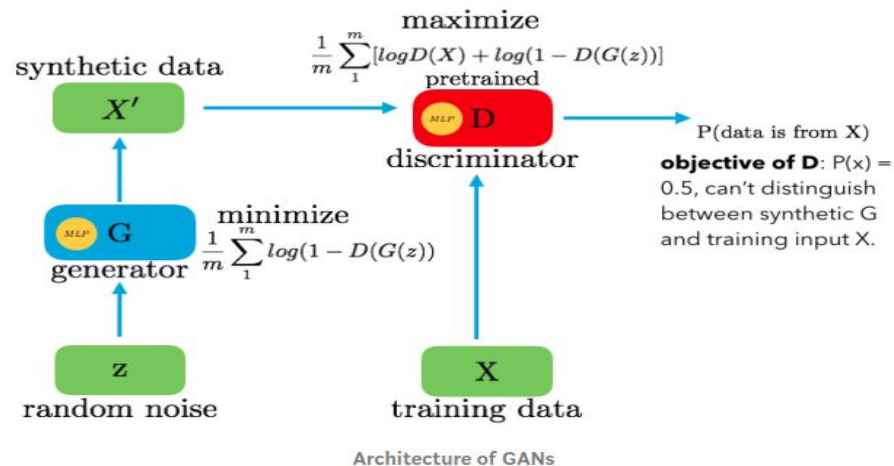
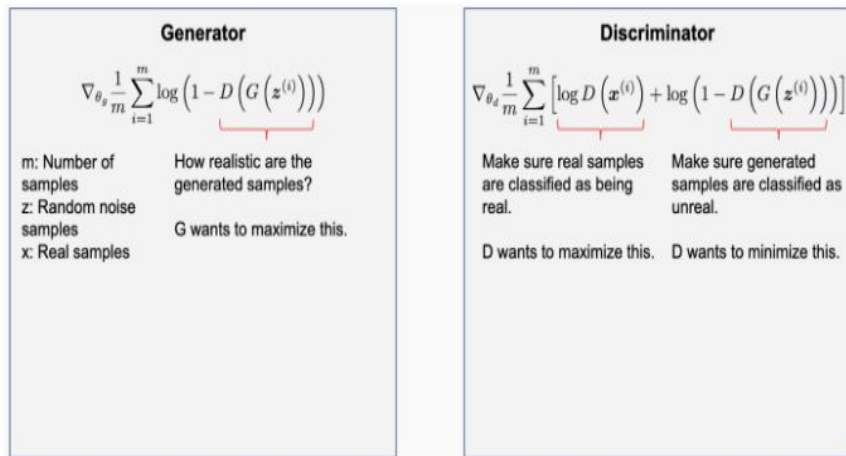
- Generator network takes as input a N-dimensional random noise and produces a fake image
- Discriminator network input is both fake samples as well as samples from the dataset
- **Generator Training Goal:** maximise D's final classification error.
  - The generated images are perceived as real
- **Discriminator Training Goal:** minimise the final classification error.
  - Real data is correctly distinguished from fake data
  - Makes binary decision about image  $G(z)$ : a real image  $D(z)=1$ , or a fake image,  $D(z)=0$ .

- Loss function should estimate the **cumulative performance** of the two networks?
- Predictions on the dataset by the discriminator should be as close to 1 as possible, and on the generator to be as close to 0 as possible.
- To achieve this, use the log-likelihood of  $D(x)$  and  $1-D(z)$  in the objective function.

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

# Loss Function: Intuition

< itc >

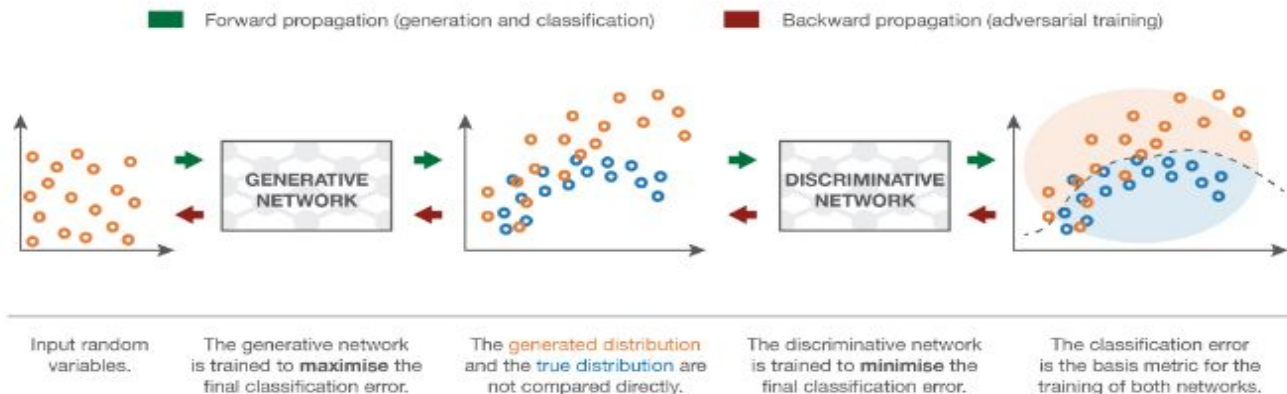


For the Generator, we want to minimize  $\log(1-D(G(z)))$  i.e. when the value of  $D(G(z))$  is high then D will assume that  $G(z)$  is nothing but X and this makes  $1-D(G(z))$  very low and we want to minimize it which this even lower. For the Discriminator, we want to maximize  $D(X)$  and  $(1-D(G(z)))$ . So the optimal state of D will be  $P(x)=0.5$ . However, we want to train the generator G such that it will produce the results for the discriminator D so that D won't be able to distinguish between z and X.

# Training Procedure: Backpropagation:

< itc >

- To train the networks we perform backpropagation whilst freezing the other network's neuron weights.
- Generator's weights are updated using Gradient Ascent to maximise the error, while Discriminator uses Gradient Descent to minimise it.



- Perform several iterations of gradient descent on D using real and generated images by fixing G
- Then fix D and train G for several iterations (usually more iterations required to train G) to fool a fixed D.
- We want to optimize our minimax function by iterating both G and D in alternating steps until discriminator won't be able to differentiate between real and fake images ( $D(x) = 0.5$ )

# Training Procedure: Algorithm



---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

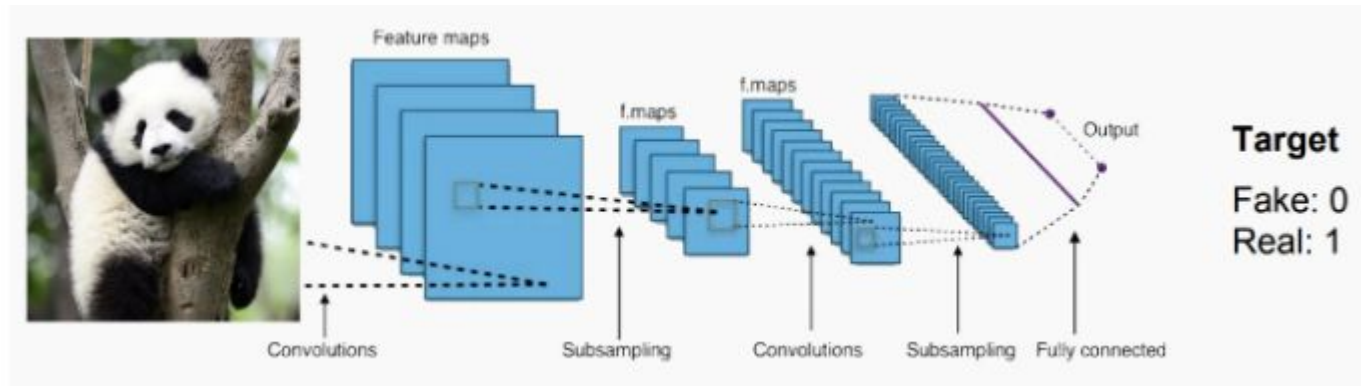
---



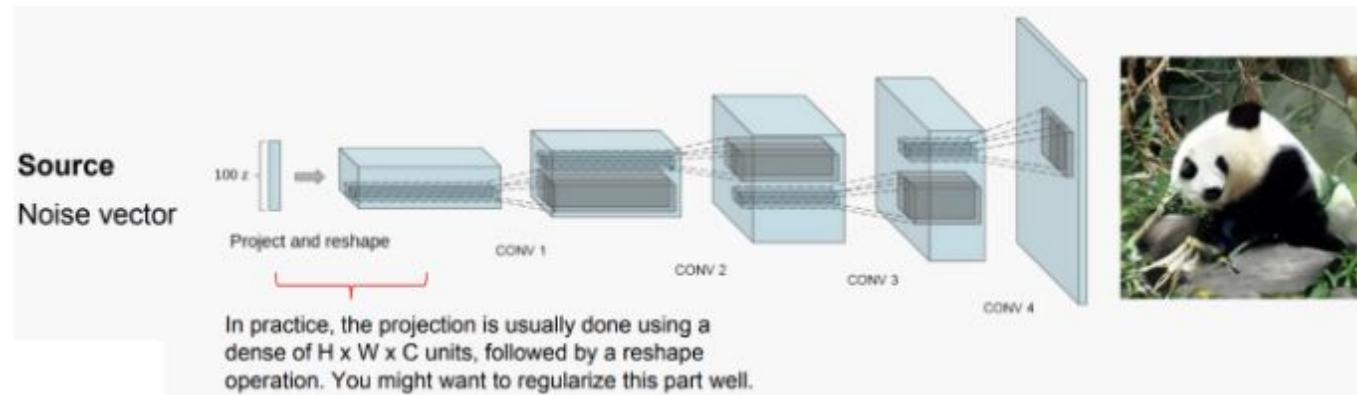
# GANs Architectures: Basics

< itc >

## Discriminator



## Generator



# GANs Architectures: Advanced Considerations (DCGANs)



- **Replace Pooling with Strided convolutions** (D learns its own spatial downsampling and G its respective upsampling without adding any bias)
- **Use BatchNorm** (stabilizes learning by normalizing the input, creates more robust deep models without having the gradients diverging)
- **Avoid using Fully-Connected hidden layers** (hurts convergence speed)
- **G: use ReLU and Tanh for the output.** (tanh is preferred activation for images as an output as it has a range of  $[-1, 1]$ )
- **D: use LeakyReLU** (tested empirically, works well for modelling to a higher resolution)

# GANs Evolution and Flavours

< itc >

GANs keeps improving!!



Reference: [The GAN Zoo on GitHub - Hindu Puravinash](#) ~ 300 GAN Papers!!!! Mostly GAN variations

- 2D-ID-GAN - Shape Inpainting using 2D Generative Adversarial Network and Recurrent Convolutional Networks
- 2D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 2D Generative-Adversarial Modeling (github)
- 2D-WGAN - Improved Adversarial Systems for 3D Object Generation and Reconstruction (github)
- 2D-RawGAN - 3D Object Reconstruction from a Single Depth View with Adversarial Learning (github)
- ABC-GAN - ABC-GAN: Adaptive Blur and Control for improved training stability of Generative Adversarial Networks (github)
- ABC-GAN - GANs for LIFE: Generative Adversarial Networks for Likelihood Free Inference
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- wGAN - Face Aging With Conditional Generative Adversarial Networks
- ACuAL - ACuAL: Actor-Critic Under Adversarial Learning
- AdaGAN - AdaGAN: Boosting Generative Models
- AdvGAN - Generating adversarial examples with adversarial networks
- AE-GAN - AE-GAN: adversarial eliminating with GAN
- WEGAN - Learning Inverse Mapping by Autoencoders Based Generative Adversarial Nets
- AIRGAN - Assorted MAP Inference for Image Super-resolution
- ALCGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference (github)
- AlignGAN - AlignGAN: Learning to Align Cross-Domain Images with Conditional Generative Adversarial Networks
- AM-GAN - Activation Maximization Generative Adversarial Nets
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- APE-GAN - APE-GAN: Adversarial Perturbation Elimination with GAN
- ARAE - Adversarially Regularized Autoencoders for Generating Discrete Structures (github)
- ARDA - Adversarial Representation Learning for Domain Adaptation
- ARIGAN - ARIGAN: Synthetic Arabidopsis Plants using Generative Adversarial Network
- ARIGAN - ARIGAN: Artwork Synthesis with Conditional Generative GANs
- ARIGAN - Arbitrary Facial Attribute Editing: Only Change What You Want
- AttGAN - AttGAN: Fine-Grained Text-to-Image Generation with Attentional Generative Adversarial Networks
- TV-GAN - TV-GAN: Generative Adversarial Network Based Thermal to Visible Face Recognition
- UGAN - Unsupervised Generative Adversarial Cross-modal Hashing
- UGAN - Enhancing Underwater Imagery using Generative Adversarial Networks
- UninGAN - Unsupervised Image-to-Image Translations with Generative Adversarial Networks (github)
- Unrolled GAN - Unrolled Generative Adversarial Networks (github)
- VAE-GAN - Autoencoding beyond pixels using a learned similarity metric
- VarGAN - Multi-View Image Generation from a Single-View
- VAW-GAN - Voice Conversion from Unaligned Corpora using Variational Autoencoding Wasserstein Generative Adversarial Networks
- VEGAN - VEGAN: Reducing Mode Collapse in GANs using Implicit Variational Learning (github)
- VGAN - Generating Videos with Scene Dynamics (github)
- VGAN - Generative Adversarial Networks as Variational Training of Energy Based Models (github)
- VGAN - Text Generation Based on Generative Adversarial Nets with Latent Variable
- VIGAN - Image Generation and Editing with Variational Info Generative Adversarial Networks
- VIGAN - VIGAN: Missing View Imputation with Generative Adversarial Networks
- VoiceGAN - Voice Impersonation using Generative Adversarial Networks
- VRAL - Variance Regularizing Adversarial Learning
- WaterGAN - WaterGAN: Unsupervised Generative Network to Enable Real-time Color Correction of Monocular Underwater Images
- WaveGAN - Synthesizing Audio with Generative Adversarial Networks
- wGAN - Generative Adversarial Nets for Multiple Text Corpora
- WGAN - Wasserstein GAN (github)
- WGAN-GP - Improved Training of Wasserstein GANs (github)
- WVGAN - Visually Supervised Generative Adversarial Networks for 3D Reconstruction
- XGAN - XGAN: Unsupervised Image-to-Image Translation for many-to-many Mappings
- ZipNet-GAN - ZipNet-GAN: Inferring Fine-grained Mobile Traffic Patterns via a Generative Adversarial Neural Network
- z-GAN - Variational Approaches for Auto-Encoding Generative Adversarial Networks (github)
- z-GAN - Triangle Generative Adversarial Networks

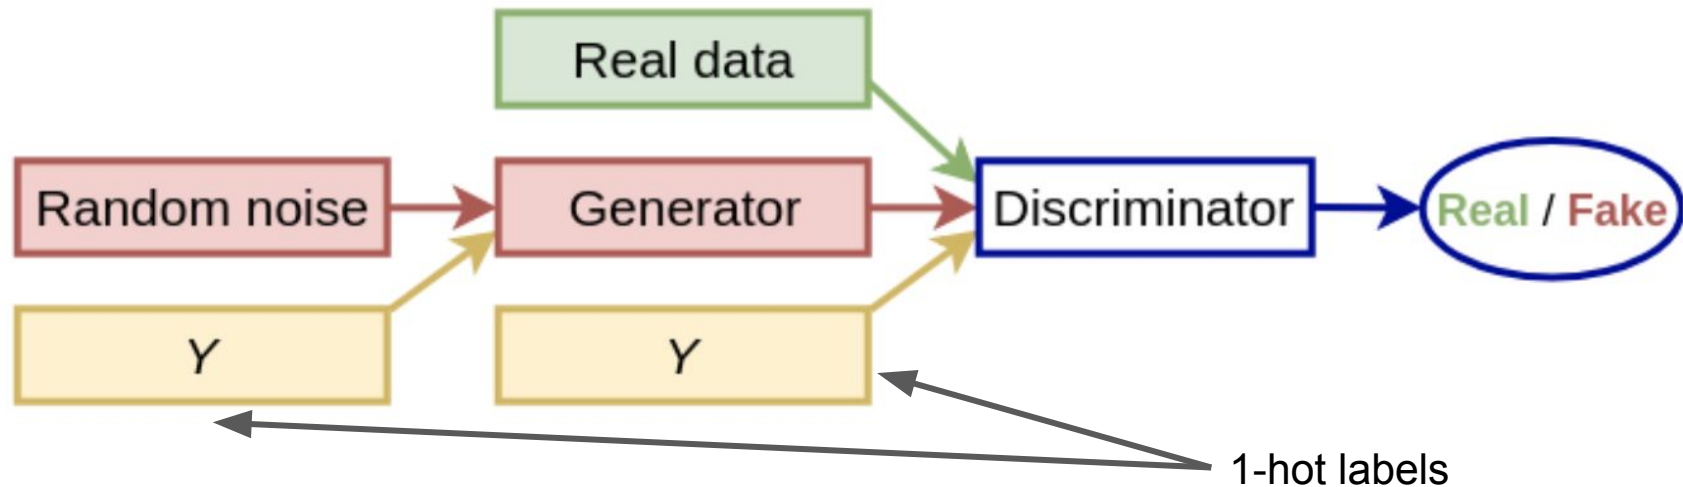
GAN's Jungle is thriving!!

- **Conditional GANs:** generate images given their features (labels)
  - Use extra label information and are able to control how generated images will look.
  - Learn to produce better images by exploiting the information fed to the model
- **Info GAN:** information-theoretic extension.
  - Learn disentangled representations in an unsupervised manner.
  - Used for very complex datasets
  - Train cGAN for unlabeled datasets to extract the most important images feature
- **Cycle GAN:** unpaired image to image translation.
  - Transforms an image from one domain (zebras images), to another (horses images)
  - Other image features( not directly related to either domain as background) stays recognizably the same

# Conditional Gans

**Discriminator:** similar to a standard Discriminator except for the 1-hot vector, which is used to condition Discriminator outputs

**Generator:** similar to a standard Generator except for the 1-hot vector, which is used to condition Generator outputs




# Concept: Disentangled vs Entangled Latent Space < itc >

- The only way to “turn” to change the generator output is the noise input
- Since it’s noise, there’s no intuition about how to modify it to get a desired effect
- What if you wanted an image of a man with glasses — how do you **change the noise**?
- **Impossible** as your representation is entangled. InfoGAN tries to solve this problem and provide a disentangled representation.



Entangled




Disentangled

# InfoGAN: meaningful latent code/consistent effects on the output


< itc >

Split Generator input into 2 parts:

- Traditional noise vector
- New “latent code” vector. The codes are then made meaningful by maximizing the [Mutual Information](#) between the code and the generator output



GAN




InfoGAN

# Cool GAN types: DaGAN(data augmentation)

< itc >

- Generates a synthetic image using a lower-dim representation of a real image(autoencoder style)
- Takes an existing image, encodes it, adds noise, and decodes it.
- Decoder learns a large family of transformations for data augmentation.



Original and transformed image (fake) vs.  
original image and different image of same class (real)



# Cool GAN types: ProGAN (progressive)

< itc >

- Grows generator and discriminator progressively:
  - Starts from a low resolution, adds new layers modeling increasingly finer details while training progresses
  - Capable to generate high-quality images compared to its predecessors


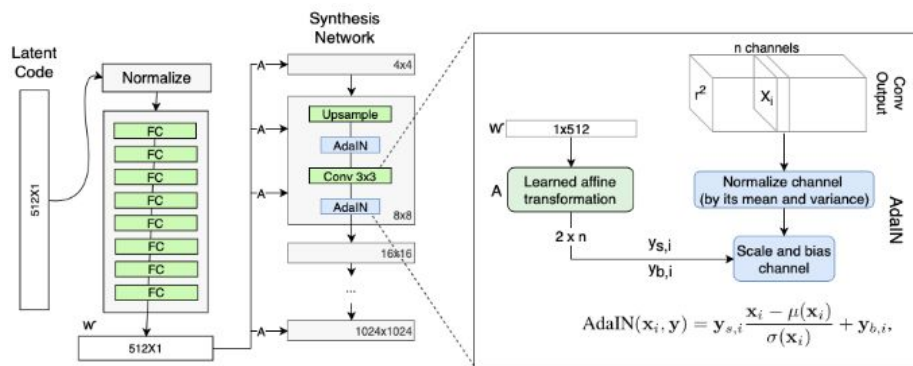


Figure 1: Our training starts with both the generator (G) and discriminator (D) having a low spatial resolution of  $4 \times 4$  pixels. As the training advances, we incrementally add layers to G and D, thus increasing the spatial resolution of the generated images. All existing layers remain trainable


# Cool GAN types: StyleGAN



- Control visual features of the generated images, such as hair color, pose, nose form etc
- Exploits ProGAN progressive layers' ability to control image visual features, if utilized properly
- The lower the layer (the resolution), the coarser the features it affects
  - Coarse - affects pose, general hair style, face shape, etc
  - Middle - affects finer facial features, hair style, eyes open/closed, etc.
  - Fine - affects color scheme (eye, hair and skin) and micro features.



The generator's Adaptive Instance Normalization (AdaIN)



- **Sensitivity** to both structure and parameters
  - If either the discriminator or generator gets better than the other too quickly, the other will never be able to catch up
  - Finding the right combination can be very challenging
- **Convergence** - No proof that GANs converge.
  - GANs perform well with the right parameters, but there's no guarantee beyond that.
  - The more complicated network gets, the tougher convergence becomes, the more difficult hyperparameter selection becomes

- **Generation of high-resolution big size images** (partially solved by ProGANs, CycleGANs but still...)
  - It's easy for “D” to tell the generated fakes from the real images
  - Many pixels can lead to error gradients that cause the Gs output to move in almost random directions
- **Mode(modal) Collapse:** “G” produces the same image every time independently of input noise (InfoGAN solves this issue partially)
  - When training our network, the “G” somehow finds one image that fools the “D”
  - “D” will always say it is real, so the “G” has accomplished its goal and stops learning.
  - However, the problem is that **every sample made by the generator is identical.**

# Further readings about GANs



- [Brief introduction to GANs](#)
- [Brief Introduction to GAN 2](#)
- [Comprehensive GAN analysis with code](#)
- [Generating dogs images with GANs\(tensorflow\)](#)
- [How to Measure GAN Performance](#)
- [GANs Flaws](#)
- [GANs type \(with some a bit advanced math\)](#)
- [StyleGAN: thorough analysis and examples](#)
- [GANs for Data Augmentation](#)

Ask the lecturer for more cool GAN Stuff